

# Probabilistic programming for mechanistic models using PyMC3

---

Sanmitra Ghosh

March 6, 2018

Computational Biology Group,  
Department of Computer Science,  
University of Oxford  
*[sanmitra.ghosh@cs.ox.ac.uk](mailto:sanmitra.ghosh@cs.ox.ac.uk)*

# Table of contents

1. Bayesian inference in mechanistic models
2. Probabilistic programming in PyMC3
3. Probabilistic Models and AutoDiff
4. Conclusion

# Bayesian inference in mechanistic models

---

# Bayesian inference summary

A generative mechanistic model:  $X \sim \mathcal{M}_\theta$ ,  $\mathcal{M}_\theta$  being an ODE system

$$\frac{dX}{dt} = f(X, \theta). \quad (1)$$

Having observed some data  $y$  we define a likelihood  $p(y|X, \theta)$  and combine that with a prior distribution  $p(\theta)$  to obtain the posterior as:

$$p(\theta|y) = \frac{p(y|X, \theta)p(\theta)}{\int p(y|X, \theta)p(\theta)d\theta}, \quad (2)$$

and the marginal likelihood or the evidence of model  $\mathcal{M}_\theta$  as:

$$p(y|\mathcal{M}_\theta) = \int p(y|X, \theta)p(\theta)d\theta. \quad (3)$$

- $p(y|\mathcal{M}_\theta)$  is intractable and thus we resort to sampling.

# Bayesian inference challenges

Conceptually MCMC is very simple. However, a modeller needs to spend some time in implementing and tuning MCMC to get useful results.


Implementing advanced MCMC schemes and/or estimating normalizing constants require further experience and training in computational statistics.

- Probabilistic programming libraries such as STAN [1] and PyMC3 [2] can make things considerably easier.

# Probabilistic programming in PyMC3

---

# Lets automate inference: Linear regression



The generative model (likelihood):

$$Y = \mathcal{N}(\mu, \sigma^2)$$
$$\mu = \alpha + \beta_1 X_1 + \beta_2 X_2 \quad (4)$$

The priors:

$$\alpha = \mathcal{N}(0, 100)$$
$$\beta_i = \mathcal{N}(0, 100) \quad (5)$$
$$\sigma = |\mathcal{N}(0, 1)|$$

# Lets automate inference: PyMC3

---

```
import pymc3 as pm
```

```
basic_model = pm.Model()
```

```
with basic_model:
```

```
# Priors for unknown model parameters
```

```
alpha = pm.Normal('alpha', mu=0, sd=10)
```

```
beta = pm.Normal('beta', mu=0, sd=10, shape=2)
```

```
sigma = pm.HalfNormal('sigma', sd=1)
```

```
# Expected value of outcome
```

```
mu = alpha + beta[0]*X1 + beta[1]*X2
```

```
# Likelihood (sampling distribution) of observations
```

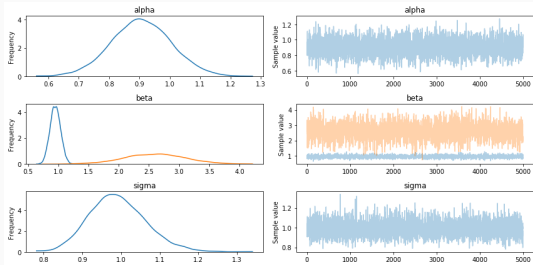
```
Y_obs = pm.Normal('Y_obs', mu=mu, sd=sigma, observed=Y)
```

---



# Lets automate inference: PyMC3...

```
with basic_model:  
    # draw 5000 posterior samples  
    trace = pm.sample()  
pm.traceplot(trace)
```




- Supported sampler: HMC [3], NUTS [4](default), Metropolis, SMC [5], Slice [6].
- PyMC3 doesn't support ODEs out-of-the-box. But there is a hack!

# Probabilistic Models and AutoDiff


---

# The PyMC3 under the hood

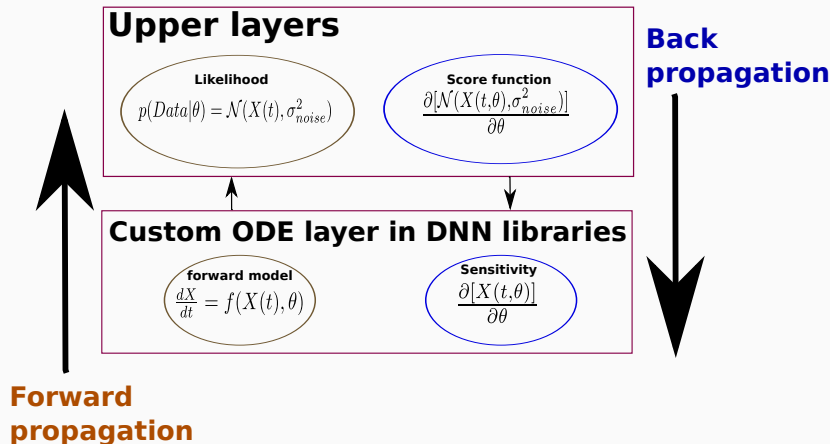
PyMC3 compiles the model into a Theano [7] program. Theano rolls out a symbolic computation graph which is compiled as a C program. Theano is used as a backend to facilitate auto-differentiation to calculate gradients of likelihood. This is required to run the Hamiltonian MC, NUTS sampler.



# ODE layer in Theano, TensorFlow etc.



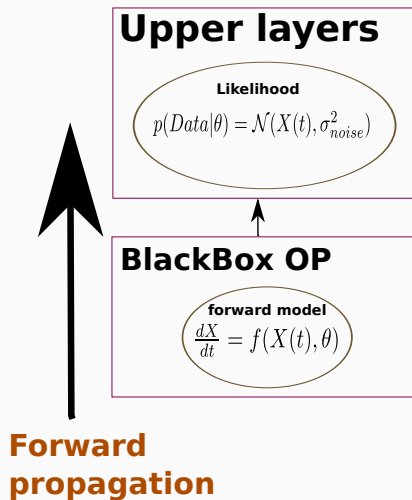
# ODE layer Backprop



- Simply apply Backprop, the Deep learning trick:


$$\frac{\partial[\mathcal{N}(X(t, \theta))]}{\partial \theta} = \frac{\partial[\mathcal{N}(X(t, \theta))]}{\partial X(t, \theta)}^T \cdot \frac{\partial X(t, \theta)}{\partial \theta} \quad (6)$$

## Or use a Black-Box layer




# Theano custom operation

We can define a custom Theano operation as a black-box wrapper over a python ODE solver. We just need to tell Theano the input,  $\theta$  and output,  $X$  data types.



- We can't use auto-differentiation and thus NUTS sampler and Variational Inference. However, we can use Metropolis and the powerful SMC sampler.

# SMC in PyMC3: Transitional Markov Chain Monte Carlo



Sample from a series of intermediate distributions to reach the target. Intermediate distributions are given by

$$\pi_t = p(\mathbf{y}|\boldsymbol{\theta})^t p(\boldsymbol{\theta}), \quad (7)$$

where  $0 < t < 1$  is used to anneal the likelihood surface.

- As a by product this algorithm returns an estimate of the marginal likelihood  $p(\mathbf{y}) = \int p(\mathbf{y}|\boldsymbol{\theta})p(\boldsymbol{\theta})$ .



# Variational inference

The Bayesian setup once again;

$$p(\theta|Data) = \frac{p(Data|\theta)p(\theta)}{p(Data)} \quad (8)$$

In variation inference we approximate  $p(\theta|Data)$  with an approximate pdf  $q(\theta|\lambda)$ , say a Normal density where  $q(\theta|\lambda) = \mathcal{N}(\theta|\lambda)$ .

- Unlike MCMC we don't sample from  $p(\theta|Data)$ , we just find the parameters  $\lambda$  that maximise a chosen distance metric  $||$  between  $q(\theta|\lambda)||p(\theta|Data)$  through gradient ascent.
- Thus our expensive sampling based inference is turned into a significantly faster deterministic optimisation.

Let me walk you through a simple example in jupyter.  
[Github: Probabilistic Programming for Mechanistic Models](#)

## Conclusion

---

# STAN vs PyMC3

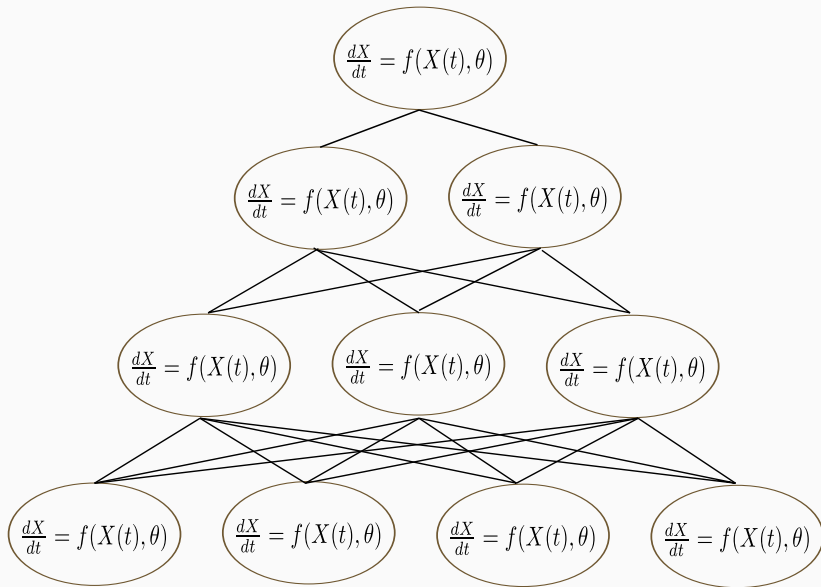
## PyMC3

- User can choose any solver. Not limited to ODEs only.
- SMC can explore the posterior surface efficiently. But computational cost is much higher.
- Probably the easiest software for model selection.
- Considerably slower than STAN.

## STAN





- One can only use the default ODE solver and NUTS sampler.
- For many well specified statistical models NUTS is the fastest sampler.
- Estimating marginal likelihood is difficult.
- Gradient based sampling is not the best idea for sloppy and/or unidentifiable models.
- Implemented in C++, so good for HPC applications.

## Aside: Deep Dynamical Net






Thank you.

# References I

-  A. Gelman, D. Lee, and J. Guo, “Stan: A probabilistic programming language for bayesian inference and optimization,” *Journal of Educational and Behavioral Statistics*, vol. 40, no. 5, pp. 530–543, 2015.
-  A. Patil, D. Huard, and C. J. Fonnesbeck, “Pymc: Bayesian stochastic modelling in python,” *Journal of statistical software*, vol. 35, no. 4, p. 1, 2010.
-  S. Duane, A. D. Kennedy, B. J. Pendleton, and D. Roweth, “Hybrid monte carlo,” *Physics letters B*, vol. 195, no. 2, pp. 216–222, 1987.
-  M. D. Hoffman and A. Gelman, “The no-u-turn sampler: adaptively setting path lengths in hamiltonian monte carlo.” *Journal of Machine Learning Research*, vol. 15, no. 1, pp. 1593–1623, 2014.

## References II

-  J. Ching and Y.-C. Chen, “Transitional markov chain monte carlo method for bayesian model updating, model class selection, and model averaging,” *Journal of engineering mechanics*, vol. 133, no. 7, pp. 816–832, 2007.
-  R. M. Neal, “Slice sampling,” *Annals of statistics*, pp. 705–741, 2003.
-  J. Bergstra, O. Breuleux, F. Bastien, P. Lamblin, R. Pascanu, G. Desjardins, J. Turian, D. Warde-Farley, and Y. Bengio, “Theano: A cpu and gpu math compiler in python,” in *Proc. 9th Python in Science Conf*, 2010, pp. 1–7.